

# 人 工 智 慧 概 論

---

113-2 week 5

# 大綱

1. SVM
2. KNN
3. 決策樹
4. 集成學習
  - 隨機森林
  - AdaBoost

# Support Vector Machine(SVM)


---

# 支持向量機 (support vector machine, SVM)

- 是一種監督式的學習方法
- 用統計風險最小化的原則來估計一個分類的超平面(hyperplane) ,
- 其基礎的概念非常簡單，就是找到一個決策邊界(decision boundary)讓兩類之間的寬度(margins)最大化，使其可以完美區隔開來。
- 分成：
  - 線性可分支持向量機
  - 非線性可分支持向量機



# 線性可分支持向量機

- 下圖是在二維圖形中找出一條線，目標讓這條直線與兩個類別之間の間隔寬度 ( margins ) 距離最大化。
- 其中離兩條虛線(間隔超平面)距離最近的點，就稱為支持向量 (support vector)。




# 非線性可分支持向量機

- 現實生活中的資料往往稍微複雜，那如果不是線性可分集合怎麼辦呢？
- 我們可以運用核函數(kernel function) 幫我們造出不可分的分割平面。




# 非線性可分支持向量機

- 除了進行線性分類之外 SVM 還可以使用核技巧有效地進行非線性分類
- 將其輸入的資料投到更高維度的空間，並在高維度的空間進行高維度的分類或降維。
- 透過多維度的投影技巧將原本在二維空間中不可分的點到了三維空間就可分了。
- 但是隨著資料量增加其運算也會變多，相對的執行速度就會變慢。



# SVM是怎麼得到那條很好的線呢？


- 直線來說，首先紅色的線會創造兩條黑色平行於紅色線的虛線，並讓黑線平行移碰到最近的一個點
- 紅線到黑線的距離稱為Margin，而SVM就是透過去找Margin最大的那個紅線，來找最好的線






# Margin最大

- 假設紅線是 $w^*x = 0$  在紅線上方的區域就是 $w^*x > 0$  紅線下方的區域就是 $w^*x < 0$



# Margin最大

同理類推來看在左邊虛線上方的區域是 $w^*x < -k$  在右邊虛線下方的區域是 $w^*x > k$ ，虛線中間不會有資料點。



# Margin最大化

- 虛線上的點 $X_1, X_2$ 就是所謂的支援向量(Support vector)
- 主要是利用支援向量來算出Margin，並最大化Margin

- 那要怎麼計算margin？


利用 $X_1$ 向量- $X_2$ 向量得到的向量投影到 $W$ 就可以了

- 接下來在 $Y^*(W^*X) \geq k$  的條件下(虛線中間沒有點)

來最大化margin

$$\begin{aligned}\text{Margin} &= \frac{W \cdot \overrightarrow{X_1 - X_2}}{2\|w\|} \\ &= \frac{2k}{2\|w\|}\end{aligned}$$

$$\text{Condition} = y^*(w^*x) \geq k$$



# 實作-鳶尾花分類

- 載入iris資料集
- 切分訓練集 (train set)與測試集 (test set)

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import os
4 from sklearn.model_selection import train_test_split
5 from sklearn import datasets
6 from sklearn.svm import SVC
7 from sklearn.metrics import accuracy_score
8
9 ### 取得當前程式放置目錄
10 current_path = os.getcwd()
11
12 ### 載入iris資料集
13 iris = datasets.load_iris()
14
15 ### 抓出資料特徵與種類標記
16 features = pd.DataFrame(iris.data, columns = iris.feature_names)
17 species = pd.DataFrame(iris.target, columns = ['species'])
18
19 ### 以70%/30%拆分訓練集與測試集
20 x_train, x_test, y_train, y_test = train_test_split(features, species, test_size = 0.3)
21
22 train_set = pd.concat([x_train, y_train], axis = 1).reset_index(drop = True)
23 test_set = pd.concat([x_test, y_test], axis = 1).reset_index(drop = True)
24
```

# 實作-鳶尾花分類

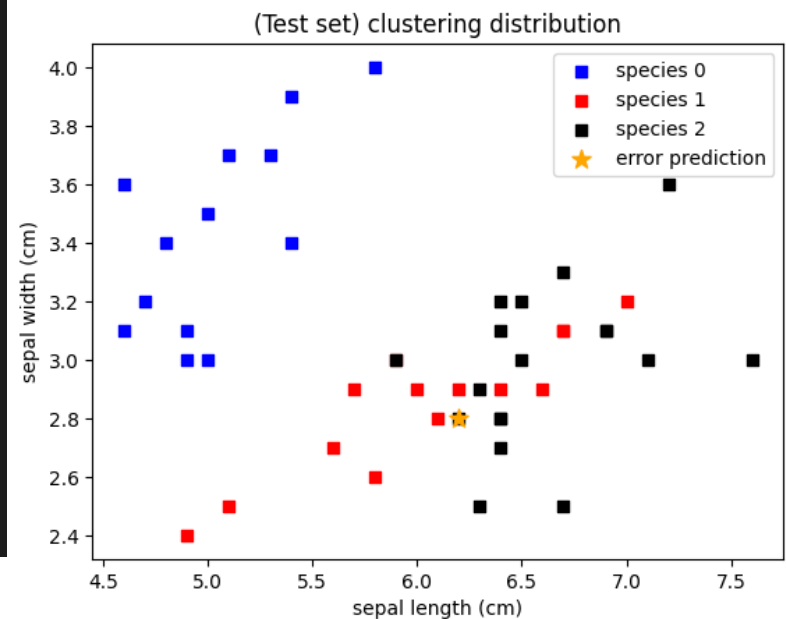
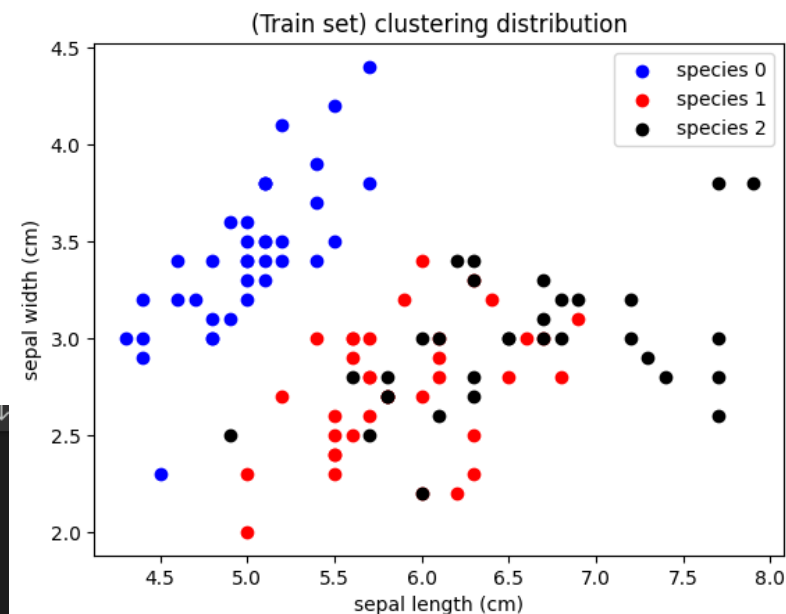
- 以訓練集的特徵（花萼、花瓣長度與寬度）、類別（花的種類）訓練SVM模型  
( 因為是分類問題，使用Support Vector Classifier, SVC )

```
25 #%% SVM訓練
26 svc = SVC()
27 svc.fit(train_set[iris.feature_names], train_set['species'])
28
29 y_pred = svc.predict(test_set[iris.feature_names])
30 test_set['species_pred'] = y_pred
31
32 # SVM分類準確率
33 SVM_accuracy = round(accuracy_score(test_set['species_pred'], test_set['species']), 4)
34 print(f'SVM accuracy: {round(SVM_accuracy * 100, 2)}%')
```

# 實作-鳶尾花分類

- 可視化結果

```
36 %% 以花瓣長度(sepal length)與花瓣寬度(sepal width)作圖
37 ### 畫出訓練集的群集分布
38 plt.figure()
39 for i in range(3):
40     if i == 0: c = 'blue'
41     if i == 1: c = 'red'
42     if i == 2: c = 'black'
43
44     item = train_set[train_set['species'] == i]
45     plt.scatter(item[iris.feature_names[0]], item[iris.feature_names[1]], color = c, label = f'species {i}')
46
47 plt.legend()
48 plt.xlabel(f'{iris.feature_names[0]}')
49 plt.ylabel(f'{iris.feature_names[1]}')
50 plt.title('(Train set) clustering distribution')
51 plt.savefig(current_path + '\\train set_clustering distribution.jpeg')
52
53 ### 畫出測試集的群集分布, 以及標記SVM預測錯誤的點
54 plt.figure()
55 for i in range(3):
56     if i == 0: c = 'blue'
57     if i == 1: c = 'red'
58     if i == 2: c = 'black'
59
60     item = test_set[test_set['species'] == i]
61     plt.scatter(item[iris.feature_names[0]], item[iris.feature_names[1]], color = c, marker = 's', label = f'species {i}')
62
63     item_pred_error = test_set[test_set['species_pred'] != test_set['species']]
64     if len(item_pred_error) != 0:
65         plt.scatter(item_pred_error[iris.feature_names[0]], item_pred_error[iris.feature_names[1]], color = 'orange', marker = '*', s = 100, label = 'error prediction')
66
67 plt.legend()
68 plt.xlabel(f'{iris.feature_names[0]}')
69 plt.ylabel(f'{iris.feature_names[1]}')
70 plt.title('(Test set) clustering distribution')
71 plt.savefig(current_path + '\\test set_clustering distribution.jpeg')
```



K Nearest Neighbor (KNN)

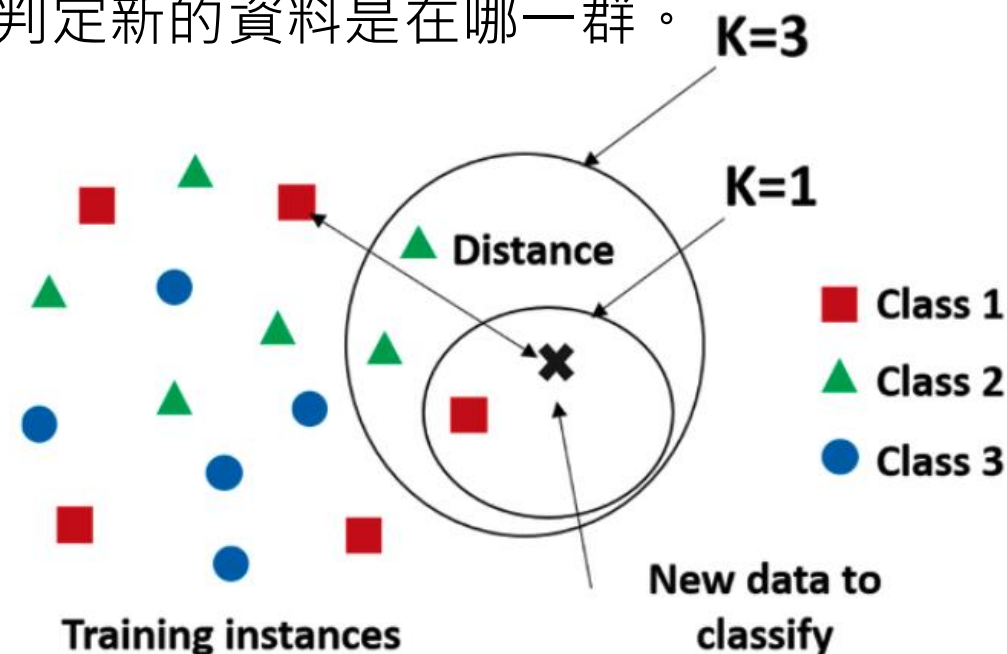
# K近鄰 ( K Nearest Neighbors, KNN )

- 一種監督式學習的分類演算法
- KNN 演算法採多數決標準，利用  $k$  個最近的鄰居來判定新的資料是在哪一群。

流程：

1. 決定  $k$  值
2. 求每個鄰居跟自己之間的距離
3. 找出跟自己最近的  $k$  個鄰居，  
查看哪一組鄰居數量最多，就加入哪一組


( 如果還是沒辦法決定在哪一組，回到第一步調整  $k$  值，再繼續 )






# K近鄰 ( K Nearest Neighbors, KNN )

- $k$  的大小會影響模型最終的分類結果。
- 以下圖為例，假設綠色點是新的資料
- 當 $k$ 等於3，會搜尋離綠色點最近的鄰居，我們可以發現藍色三角形為預測的結果。
- 當 $k$ 設為5，結果又不一样了，我們發現距離最近的三個鄰居為紅色正方形。




K=3 分類為 ▲



K=5 分類為 ■

# KNN 迴歸器

- KNN 同時也能運用在迴歸問題上面。
- 迴歸模型輸出的結果是一個連續性數值，預測該值是  $k$  個最近鄰居輸出的平均值。
- 以下圖為例當  $k=2$  時，假設我們有一個輸入特徵  $x$  要預測的輸出為  $y$ 。
- 當有一筆新的  $x$  進來的時候，KNN 迴歸器會尋找鄰近 2 個  $x$  的輸出做平均當作是該筆資料的預測結果。




# KNN 度量距離的方法


| identifier    | class name          | args    | distance function                |
|---------------|---------------------|---------|----------------------------------|
| "euclidean"   | EuclideanDistance   | •       | $\sqrt{\sum (x - y)^2}$          |
| "manhattan"   | ManhattanDistance   | •       | $\sum  x - y $                   |
| "chebyshev"   | ChebyshevDistance   | •       | $\max( x - y )$                  |
| "minkowski"   | MinkowskiDistance   | p       | $\sum  x - y ^p)^{1/p}$          |
| "wminkowski"  | WMinkowskiDistance  | p, w    | $\sum  w * (x - y) ^p)^{1/p}$    |
| "seuclidean"  | SEuclideanDistance  | V       | $\sqrt{\sum (x - y)^2 / V)}$     |
| "mahalanobis" | MahalanobisDistance | V or VI | $\sqrt{(x - y)' V^{-1} (x - y)}$ |

# KNN 與 k-means 勿混淆

- KNN 的缺點是對資料的局部結構非常敏感，因此調整適當的  $k$  值極為重要
- 大家很常將 KNN 與 K-means 混淆，雖然兩者都有  $k$  值要設定但其實兩者無任何關聯。
- KNN 的  $k$  是設定鄰居的數量採多數決作為輸出的依據。
- K-means 的  $k$  是設定集群的類別中心點數量。



K-means



KNN


# 決 策 樹

---

# Decision Trees

- Efficient method for producing **classifiers** from data - *(logic method)*.
- **Supervised learning methods** that construct decision trees from a set of input-output samples.
- It guarantees that a **simple**, but not necessary the simplest, tree is found.


- A decision tree consists of **nodes** that are **tests** on the attributes.
- The outgoing **branches** of a node correspond to all the possible **outcomes** of the test at the node.
- **Leafs** are sets of samples belong to the same class



# Decision Trees

## Credit Analysis

| salary | education      | label  |
|--------|----------------|--------|
| 10000  | high school    | reject |
| 40000  | under graduate | accept |
| 15000  | under graduate | reject |
| 75000  | graduate       | accept |
| 18000  | graduate       | accept |



# Classification by Decision Tree Induction

1. Decision tree **generation** consists of two phases:


- *Tree construction:*
  - At start, all the training examples are at the root.
  - Partition examples recursively based on selected attributes.
- *Tree pruning:*
  - Identify and remove branches that reflect noise or outliers.

2. **Use** of decision tree: Classifying an unknown sample



- Test the attribute values of the sample against the decision tree




# Decision Tree: Example




# Decision Tree: Example



# Decision Tree: Example



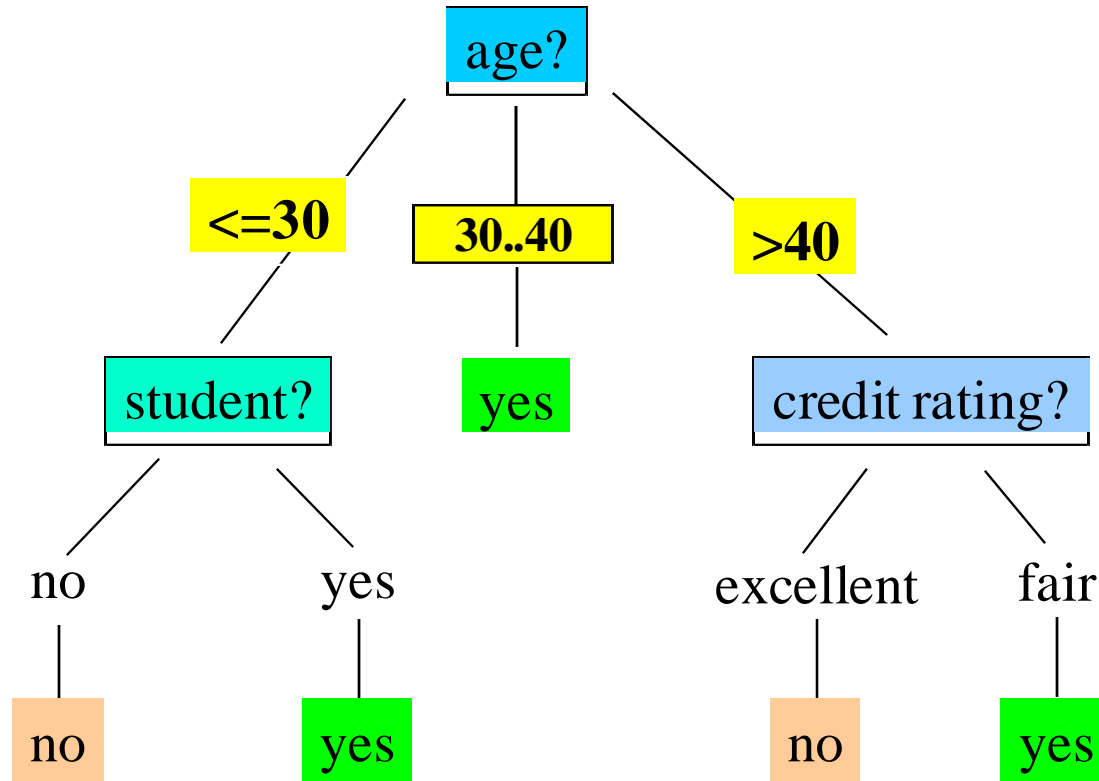
Note: tree boundaries are piecewise linear and axis-parallel



# Training Dataset

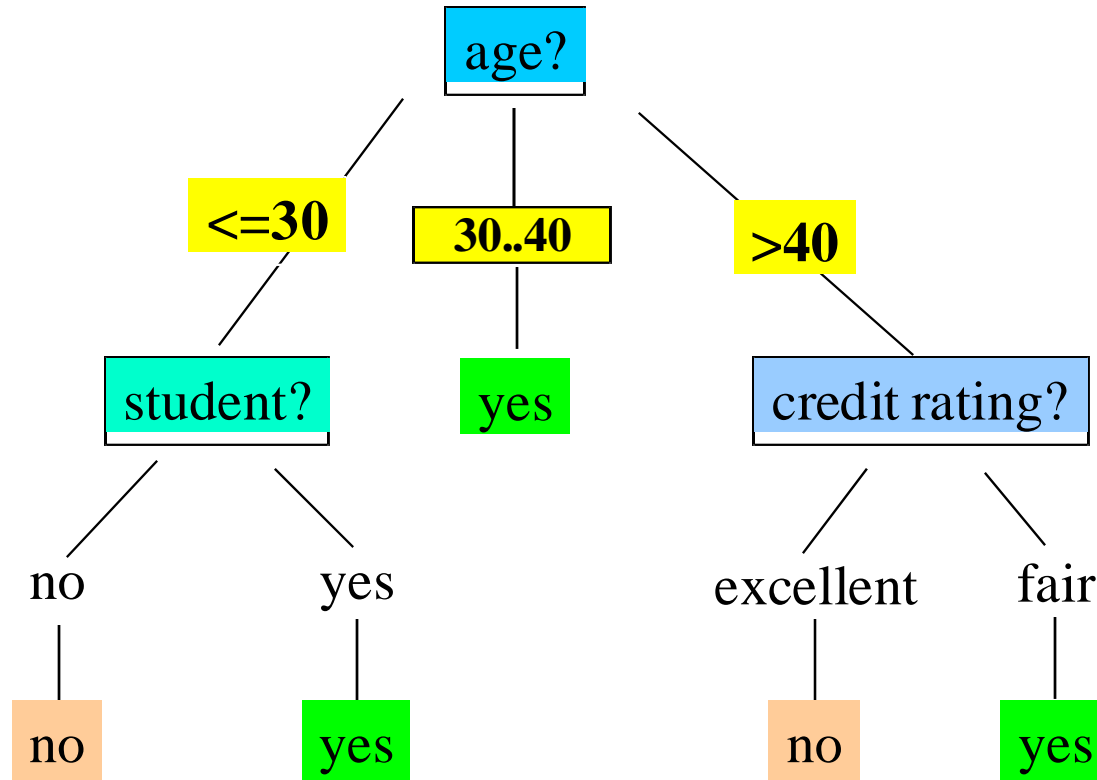
| age         | income | student | credit_rating | Buy-computer |
|-------------|--------|---------|---------------|--------------|
| < = 3 0     | high   | no      | fair          | No           |
| < = 3 0     | high   | no      | excellent     | No           |
| 3 1 ... 4 0 | high   | no      | fair          | Yes          |
| > 4 0       | medium | no      | fair          | Yes          |
| > 4 0       | low    | yes     | fair          | Yes          |
| > 4 0       | low    | yes     | excellent     | No           |
| 3 1 ... 4 0 | low    | yes     | excellent     | Yes          |
| < = 3 0     | medium | no      | fair          | No           |
| < = 3 0     | low    | yes     | fair          | Yes          |
| > 4 0       | medium | yes     | fair          | Yes          |
| < = 3 0     | medium | yes     | excellent     | Yes          |
| 3 1 ... 4 0 | medium | no      | excellent     | Yes          |
| 3 1 ... 4 0 | high   | yes     | fair          | Yes          |
| > 4 0       | medium | no      | excellent     | No           |

# Training Output: A Decision Tree for “buys\_computer”



| age     | income | student | credit_rating |     |
|---------|--------|---------|---------------|-----|
| <=30    | high   | no      | fair          | No  |
| <=30    | high   | no      | excellent     | No  |
| 31...40 | high   | no      | fair          | Yes |
| >40     | medium | no      | fair          | Yes |
| >40     | low    | yes     | fair          | Yes |
| >40     | low    | yes     | excellent     | No  |
| 31...40 | low    | yes     | excellent     | Yes |
| <=30    | medium | no      | fair          | No  |
| <=30    | low    | yes     | fair          | Yes |
| >40     | medium | yes     | fair          | Yes |
| <=30    | medium | yes     | excellent     | Yes |
| 31...40 | medium | no      | excellent     | Yes |
| 31...40 | high   | yes     | fair          | Yes |
| >40     | medium | no      | excellent     | No  |

## Testing A Decision Tree for “buys\_computer”



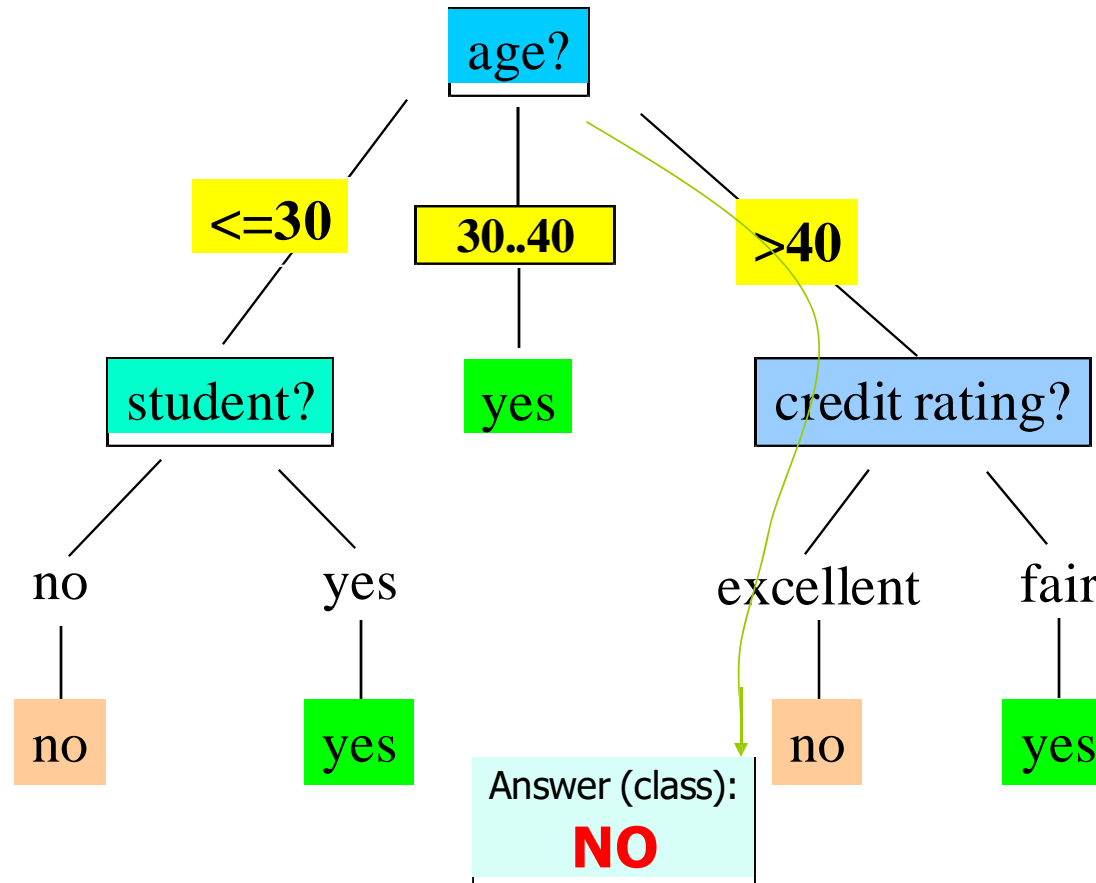
**Classify new sample:**

{ **Age = 45** and **income = low** and **student = no**  
and **credit = excellent** }

=====

Buy computer = ?

## Output: A Decision Tree for “buys\_computer”



**Classify new sample:**

{ **Age = 45** and **income = low** and **student = no** and **credit = excellent** }

=====

Buy computer = ?

# Decision Tree Algorithms

## Requirements:

1. Attribute-value description
2. Predefined classes
3. Discrete classes
4. Sufficient data
5. "Logical" classification models



# Decision Trees

## Pros

- Fast execution time
- Generated trees (rules) are easy to interpret by humans
- Scale well for large data sets
- Can handle high dimensional data

## • Cons

- Cannot capture correlations among attributes
- Consider only axis-parallel cuts

# Decision Tree Algorithms

- Classifiers from machine learning and statistical community:
  - ID3
  - C4.5[Quinlan 93] → C5.0
  - CART

# Algorithm for Decision Tree Induction

- **Basic algorithm** (*a greedy algorithm*):
  - Tree is constructed in a **top-down recursive divide-and-conquer manner**.
  - At start, all the training examples are at the root.
  - Attributes are categorical (if continuous-valued, they are discretized in advance).
  - Examples are partitioned recursively based on selected attributes.
  - Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**).
- **Conditions for stopping partitioning:**
  - All samples for a given node belong to **the same class**.
  - There are no remaining attributes for further partitioning – **majority voting** is employed for classifying the leaf.
  - There are **no samples** left.

# Attribute Selection Measure

- **Information gain** (ID3/C4.5)
  - All attributes are assumed to be categorical.
  - Can be modified for continuous-valued attributes.
- **Gini index** (IBM IntelligentMiner, CART, ...)
  - All attributes are assumed continuous-valued.
  - Assume there exist several possible split values for each attribute.
  - Can be modified for categorical attributes.

ID3

# 資訊理論 (information theory)

- 假設一個事件有  $n$  種結果，發生的機率分別為  $P(v_1), \dots, P(v_n)$ ，這些機率都是已知的，則定義這個事件發生後所得到的資訊量為：

$$I(P(v_1), \dots, P(v_n)) = \sum_{i=1}^n -P(v_i) \log_2 P(v_i)$$

- 各種結果發生機率愈平均，所求資訊量也愈大
- 資訊量可以當作亂度 (Entropy) 的指標，資訊量愈大，表示亂度愈大
- 解決屬性選擇的問題

# 亂度(Entropy)

- 昆蘭 (Quinlan)(1979)提出，以雪南 (Shannon) (1949)的資訊理論 (Information theory)為依據。
- 資訊理論:若一事件有 $k$ 種結果,對應的機率為 $P_i$ 。則此事件發生後所得到的資訊量 $I$ (視為Entropy)為：

$$I=-(p_1*\log_2(p_1)+ p_2*\log_2(p_2)+...+ p_k*\log_2(p_k))$$

- Example 1: 設  $k=4 \rightarrow p_1=0.25, p_2=0.25, p_3=0.25, p_4=0.25$   
 $I=-(.25*\log_2(.25)*4)=2$
- Example 2: 設  $k=4 \rightarrow p_1=0, p_2=0.5, p_3=0, p_4=0.5$   
 $I=-(.5*\log_2(.5)*2)=1$
- Example 3: 設  $k=4 \rightarrow p_1=1, p_2=0, p_3=0, p_4=0$   
 $I=-(1*\log_2(1))=0$
- 用熵(Entropy) 衡量資料的一致性

# ID3 演算法(C4.5,C5.0)

- 資訊獲利(Information Gain)

若分類標記(Y)分為(成功、失敗)兩種，X為預測變項(類別屬性;k類)，n為總樣本數( $n_1$ 為總樣本數中具成功標記的個數)，經由X變項將樣本分類後 $m_i$ 為X=i類中的總樣本個數( $m_{i1}$ 為X=i類中具成功標記的個數)。根據變項X將n個樣本分為 $m_1, m_2, \dots, m_k$ 的資訊獲利為:

$$Gain(X) = I(n, n_1) - E(X),$$

其中

$$I(n, n_1) = -((n_1/n) \log_2(n_1/n) + (1 - n_1/n) \log_2(1 - n_1/n))$$

$$E(X) = (m_1/n) * I(m_1, m_{11}) + (m_2/n) * I(m_2, m_{21}) + \dots + (m_k/n) * I(m_k, m_{k1})$$



# 利用資訊獲利做屬性選取

- 資訊獲利 =
- 測試前的資訊量 - 測試後的資訊量
- 決策樹分類的目的
  - 訓練樣本分成亂度最小的子集合
  - 所有樣本都屬於同一分類標記的子集合
- ID3
  - 以測試後資訊量最小的屬性為優先選取
  - 也就是選擇資訊獲利最大的屬性

# Example(Information Gain)

$n=16$


$n_1=4$

$$I(16,4)=-((4/16)*\log_2(4/16)+(12/16)*\log_2(12/16))=0.8113$$


$$E(\text{年齡})=(6/16)*I(6,1)+(10/16)*I(10,3)=0.7946$$

$$\text{Gain}(\text{年齡})=I(16,4)-E(\text{年齡})=0.0167$$


| 年齡  | 性別     | 家庭所得 | 購買RV房車 |
|-----|--------|------|--------|
| <35 | Male   | 小康   | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Female | 低所得  | 否      |
| <35 | Male   | 高所得  | 否      |
| ≥35 | Female | 低所得  | 否      |
| <35 | Female | 低所得  | 否      |
| <35 | Female | 高所得  | 是      |
| ≥35 | Male   | 小康   | 是      |
| <35 | Male   | 高所得  | 否      |
| ≥35 | Female | 小康   | 否      |
| <35 | Male   | 低所得  | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Male   | 低所得  | 是      |
| ≥35 | Male   | 小康   | 是      |
| ≥35 | Female | 低所得  | 否      |



$$\text{Gain}(\text{年齡})=0.0167$$



$$\text{Gain}(\text{性別})=0.0972$$




$$\text{Gain}(\text{家庭所得})=0.0177$$


Max:作為第一個分類依據

| 年齡  | 性別   | 家庭所得 | 購買RV房車 |
|-----|------|------|--------|
| <35 | Male | 小康   | 否      |
| <35 | Male | 低所得  | 否      |
| <35 | Male | 高所得  | 否      |
| <35 | Male | 高所得  | 否      |
| ≥35 | Male | 小康   | 是      |
| ≥35 | Male | 小康   | 是      |
| ≥35 | Male | 低所得  | 是      |


| 年齡  | 性別     | 家庭所得 | 購買RV房車 |
|-----|--------|------|--------|
| <35 | Female | 低所得  | 否      |
| <35 | Female | 高所得  | 是      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Female | 低所得  | 否      |
| ≥35 | Female | 低所得  | 否      |
| ≥35 | Female | 低所得  | 否      |



$$I(7,3) = -((3/7) * \log_2(3/7) + (4/7) * \log_2(4/7)) = 0.9852$$




■ Gain(年齡)=0.9852




■ Gain(家庭所得)=0.688

$$I(9,1) = -((1/9) * \log_2(1/9) + (8/9) * \log_2(8/9)) = 0.5032$$



■ Gain(年齡)=0.2222




■ Gain(家庭所得)=0.5032

# Example(end) ID3演算法

## ■ 資料

| 年齡  | 性別     | 家庭所得 | 購買RV房車 |
|-----|--------|------|--------|
| <35 | Male   | 小康   | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Female | 低所得  | 否      |
| <35 | Male   | 高所得  | 否      |
| ≥35 | Female | 低所得  | 否      |
| <35 | Female | 低所得  | 否      |
| <35 | Female | 高所得  | 是      |
| ≥35 | Male   | 小康   | 是      |
| <35 | Male   | 高所得  | 否      |
| ≥35 | Female | 小康   | 否      |
| <35 | Male   | 低所得  | 否      |
| ≥35 | Female | 小康   | 否      |
| ≥35 | Male   | 低所得  | 是      |
| ≥35 | Male   | 小康   | 是      |
| ≥35 | Female | 低所得  | 否      |

## ■ Decision Tree



## 分類規則

IF 性別=Female AND 家庭所得=低所得 THEN 購買RV房車=否

IF 性別=Female AND 家庭所得=小康 THEN 購買RV房車=否

IF 性別=Female AND 家庭所得=高所得 THEN 購買RV房車=是

IF 性別=Male AND 年齡<35 THEN 購買RV房車=否

IF 性別=Male AND 年齡≥35 THEN 購買RV房車=是

# ID3的問題

- ID3演算法所使用的資訊獲利

- 會傾向選擇擁有許多不同數值的屬性
- 例如：“產品編號”欄位中，每一個產品的產品編號皆不同。
- 依產品編號進行分割，會產生出許多分支，且每一個分支都是很單一的結果，其資訊獲利會最大。但這個屬性對於建立決策樹是沒有意義的。

- C4.5演算法

- 利用屬性的獲利比率(Gain Ratio)
- 資訊獲利正規化
- 擁有最大獲利比例的屬性被設為分割屬性

# Decision Tree Algorithms – C4.5

- **Recursive Building Tree Phase:**
  - *Initialize root node of tree.*
  - ***while** a node  $N$  that can be split:*
    - for each** attribute  $A$ , evaluate splits on  $A$ , use best split to split  $N$ .*
- Use *Entropy index* to find best split.
- Separate attribute lists maintained in each node of tree.

## C4.5 Algorithm

Let the classes be denoted  $\{C_1, C_2, \dots, C_k\}$ . There are three possibilities for the content of the set of training samples  $T$  in the given node of decision tree:

1.  *$T$  contains one or more samples, all belonging to a single class  $C_j$ :* The decision tree for  $T$  is a leaf identifying class  $C_j$ .
2.  *$T$  contains no samples:*

The decision tree is again a leaf, (the class to be associated with the leaf must be determined from information other than  $T$ )

C4.5 algorithm uses as a criterion the most frequent class at the **parent** of the given node.

## C4.5 Algorithm

### 3. *T contains samples that belong to a mixture of classes:*

In this situation, the idea is to refine  $T$  into subsets of samples that are heading towards single-class collections of samples. An **appropriate test** is chosen, based on single attribute, that has one or more mutually exclusive outcomes  $\{O_1, O_2, \dots, O_n\}$ :

- $T$  is partitioned into subsets  $T_1, T_2, \dots, T_n$  where  $T_i$  contains all the samples in  $T$  that have outcome  $O_i$  of the chosen test.
- The decision tree for  $T$  consists of a decision node identifying the test and one branch for each possible outcome.



# C4.5 Algorithm

## *An Example*

*Database T:*

| Attribute1 | Attribute2 | Attribute3 | Class  |
|------------|------------|------------|--------|
| A          | 70         | True       | CLASS1 |
| A          | 90         | True       | CLASS2 |
| A          | 85         | False      | CLASS2 |
| A          | 95         | False      | CLASS2 |
| A          | 70         | False      | CLASS1 |
| B          | 90         | True       | CLASS1 |
| B          | 78         | False      | CLASS1 |
| B          | 65         | True       | CLASS1 |
| B          | 75         | False      | CLASS1 |
| C          | 80         | True       | CLASS2 |
| C          | 70         | True       | CLASS2 |
| C          | 80         | False      | CLASS1 |
| C          | 80         | False      | CLASS1 |
| C          | 96         | False      | CLASS1 |

# C4.5 Algorithm

Database T: (An Example)

| Attribute1 | Attribute2 | Attribute3 | Class  |
|------------|------------|------------|--------|
| A          | 70         | True       | CLASS1 |
| A          | 90         | True       | CLASS2 |
| A          | 85         | False      | CLASS2 |
| A          | 95         | False      | CLASS2 |
| A          | 70         | False      | CLASS1 |
| B          | 90         | True       | CLASS1 |
| B          | 78         | False      | CLASS1 |
| B          | 65         | True       | CLASS1 |
| B          | 75         | False      | CLASS1 |
| C          | 80         | True       | CLASS2 |
| C          | 70         | True       | CLASS2 |
| C          | 80         | False      | CLASS1 |
| C          | 80         | False      | CLASS1 |
| C          | 96         | False      | CLASS1 |



$$\text{Info}(T) = -9/14 \log_2(9/14) - 5/14 \log_2(5/14) = 0.940 \text{ bits}$$

Attribute1:

=====

$$\begin{aligned} \text{Info}_{x_1}(T) &= 5/14 (-2/5 \log_2(2/5) - 3/5 \log_2(3/5)) \\ &\quad + 4/14 (-4/4 \log_2(4/4) - 0/4 \log_2(0/4)) \\ &\quad + 5/14 (-3/5 \log_2(3/5) - 2/5 \log_2(2/5)) \\ &= 0.694 \text{ bits} \end{aligned}$$

$$\text{Gain}(x_1) = 0.940 - 0.694 = \mathbf{0.246 \text{ bits}}$$

Attribute3:

=====

$$\begin{aligned} \text{Info}_{x_2}(T) &= 6/14 (-3/6 \log_2(3/6) - 3/6 \log_2(3/6)) \\ &\quad + 8/14 (-6/8 \log_2(6/8) - 2/8 \log_2(2/8)) \\ &= 0.892 \text{ bits} \end{aligned}$$

$$\text{Gain}(x_2) = 0.940 - 0.892 = \mathbf{0.048 \text{ bits}}$$

## C4.5 Algorithm

- b) If attribute  $Y$  has **continuous numeric values**, a binary test with outcomes  $Y \leq Z$  and  $Y > Z$  could be defined, based on comparing the value of attribute against a threshold value  $Z$ .

### Threshold value $Z$ :

- The training samples are first sorted on the values of the attribute  $Y$  being considered. There are only a finite number of these values, so let us denote them in sorted order as  $\{v_1, v_2, \dots, v_m\}$ .
- Any threshold value lying between  $v_i$  and  $v_{i+1}$  will have the same effect of dividing the cases into those whose value of the attribute  $Y$  lies in  $\{v_1, v_2, \dots, v_i\}$  and those whose value is in  $\{v_{i+1}, v_{i+2}, \dots, v_m\}$ . There are thus only  **$m-1$  possible splits on  $Y$** , all of which should be examined systematically to obtain an optimal split.
- It is usual to choose the midpoint of each interval:  $(v_i + v_{i+1})/2$  as the representative threshold. C4.5 chooses as the threshold a **smaller value  $v_i$**  for every interval  $\{v_i, v_{i+1}\}$ , rather than the midpoint itself.

## C4.5 Algorithm

Database T: **An Example**

| Attribute1 | Attribute2 | Attribute3 | Class  |
|------------|------------|------------|--------|
| A          | 70         | True       | CLASS1 |
| A          | 90         | True       | CLASS2 |
| A          | 85         | False      | CLASS2 |
| A          | 95         | False      | CLASS2 |
| A          | 70         | False      | CLASS1 |
| B          | 90         | True       | CLASS1 |
| B          | 78         | False      | CLASS1 |
| B          | 65         | True       | CLASS1 |
| B          | 75         | False      | CLASS1 |
| C          | 80         | True       | CLASS2 |
| C          | 70         | True       | CLASS2 |
| C          | 80         | False      | CLASS1 |
| C          | 80         | False      | CLASS1 |
| C          | 96         | False      | CLASS1 |

**Attribute2:** - After a sorting process, the set of values is:  
{65, 70, 75, 78, 80, 85, 90, 95, 96},  
and the set of potential threshold values Z is:

{65, 70, 75, 78, 80, 85, 90, 95}.

- The optimal Z value is  $Z=80$ , and the corresponding process of information gain computation for the test  $x_3$  ( $\text{Attribute2} \leq 80$  or  $\text{Attribute2} > 80$ )

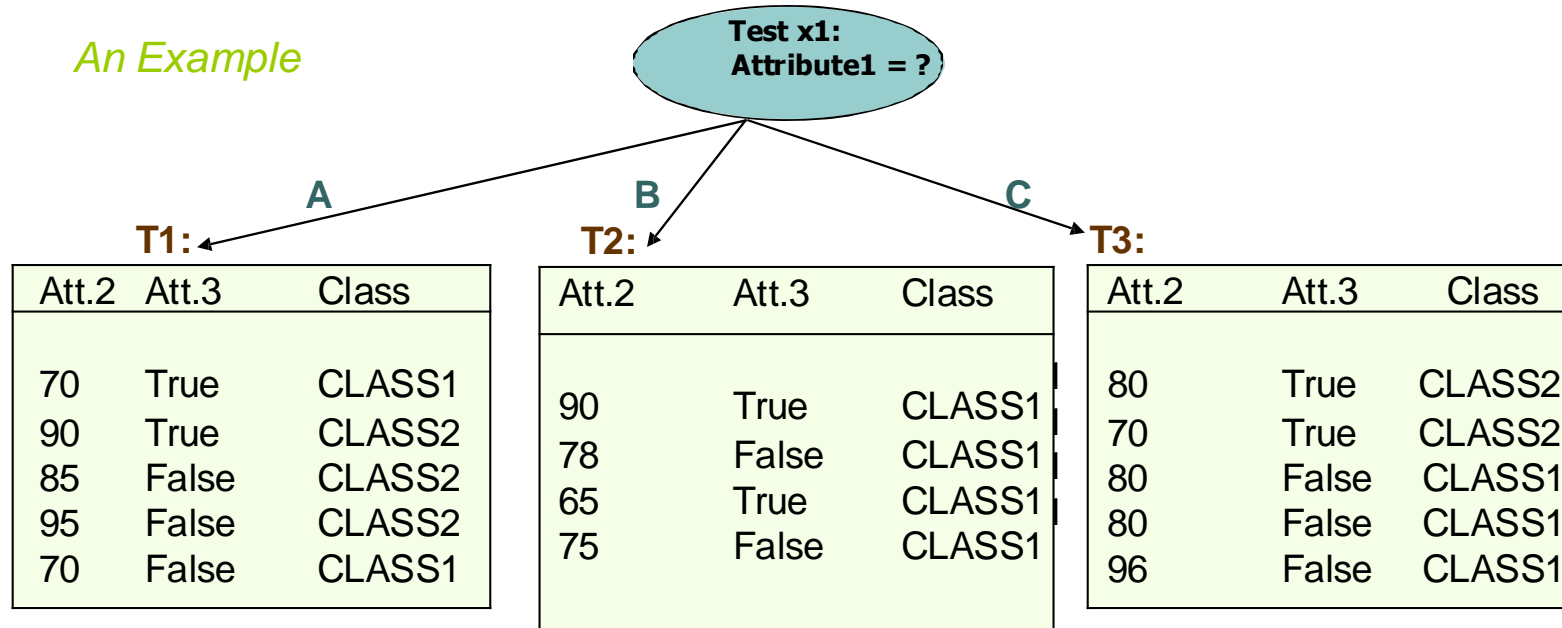
$$\text{Info}_{x_3}(T) = \frac{9}{14} (-\frac{7}{9} \log_2 (\frac{7}{9}) - \frac{2}{9} \log_2 (\frac{2}{9})) + \frac{5}{14} (-\frac{2}{5} \log_2 (\frac{2}{5}) - \frac{3}{5} \log_2 (\frac{3}{5})) = 0.837 \text{ bits}$$

$$\text{Gain}(x_3) = 0.940 - 0.837 = \mathbf{0.103 \text{ bits}}$$

**Attribute1** gives the highest gain of **0.246 bits**, and therefore this attribute will be selected for the first splitting.

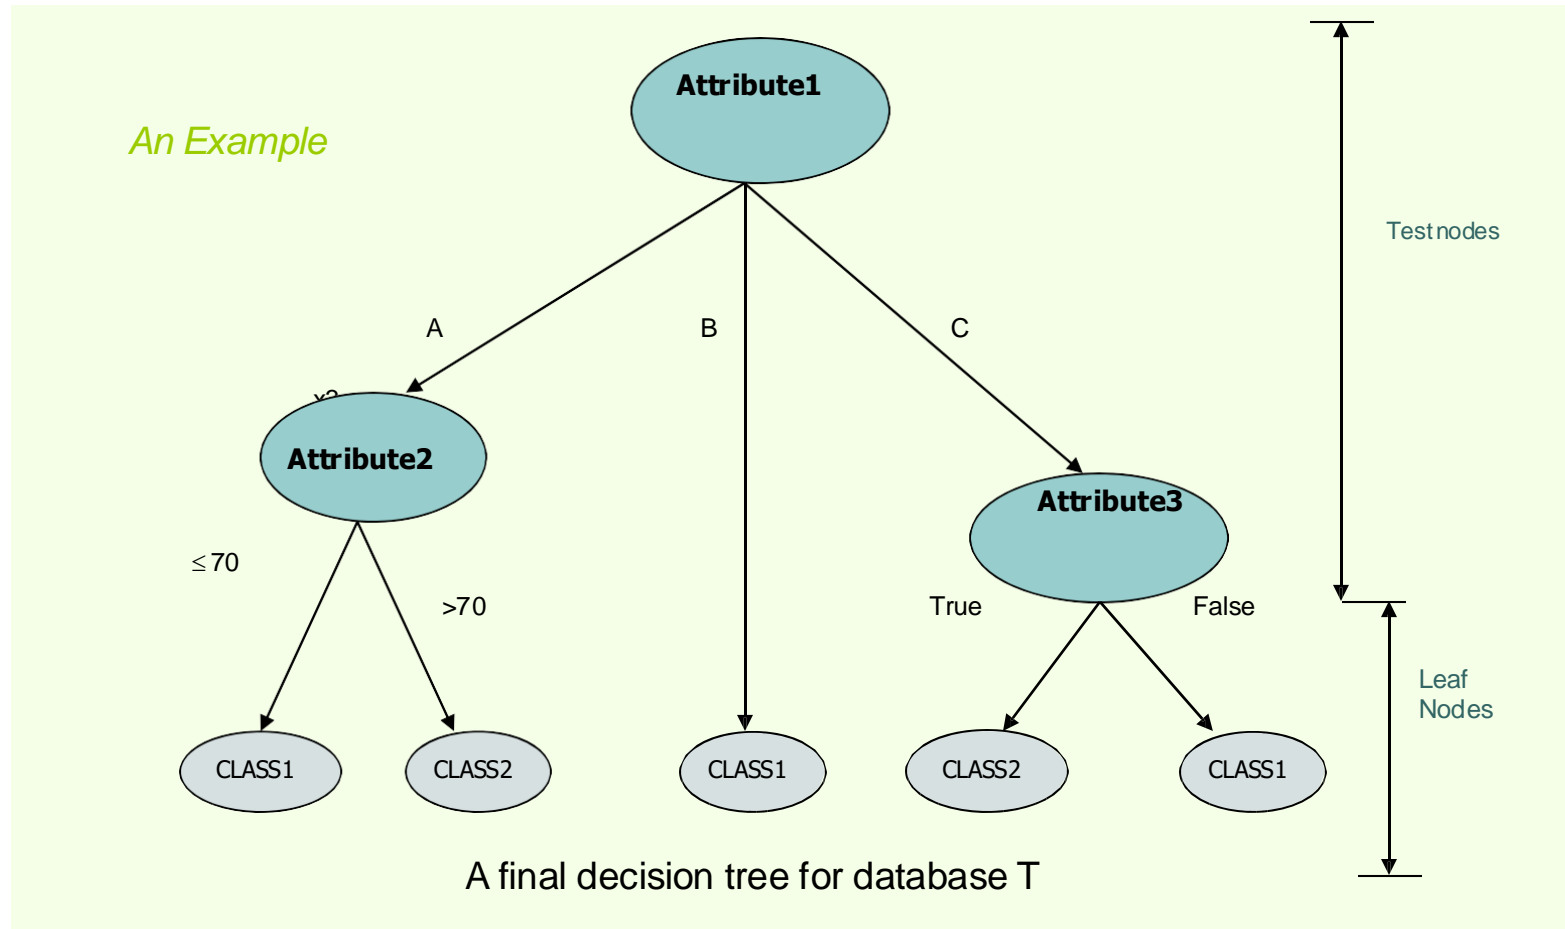
# C4.5 Algorithm

*An Example*



Initial decision tree and subset cases for a database **T**

# C4.5 Algorithm



# C4.5 Algorithm

## Decision Tree – Pseudo-code:

*An Example*

```
If      Attribute1 = A
  Then
    If      Attribute2 <= 70
      Then
        Classification = CLASS1;
      Else
        Classification = CLASS2;
    Elseif  Attribute1 = B
      Then
        Classification = CLASS1;
    Elseif  Attribute1 = C
      Then
        If      Attribute3 = True
          Then
            Classification = CLASS2;
          Else
            Classification = CLASS1.
```

## C4.5 Algorithm

- Gain criterion has **a serious deficiency**: a strong bias in favor of tests with many outcomes.
- Solution was find in some kind of normalization:

$$\text{Split-info}(X) = - \sum_{i=1}^n ((T_i / |T|) \log_2 (T_i / |T|))$$

$$\text{Gain-ratio}(X) = \text{Gain}(X) / \text{Split-info}(X)$$

Example:

$$\text{Split-info}(x_1) = - 5/14 \log_2 (5/14) - 4/14 \log_2 (4/14) - 5/14 \log_2 (5/14) = 1.577 \text{ bits}$$

$$\text{ratio}(x_1) = 0.246 / 1.557 = \mathbf{0.156}$$



# Decision Tree Algorithms

- **Building phase**

- Recursively split nodes using best splitting attribute for node.

- **Pruning phase**

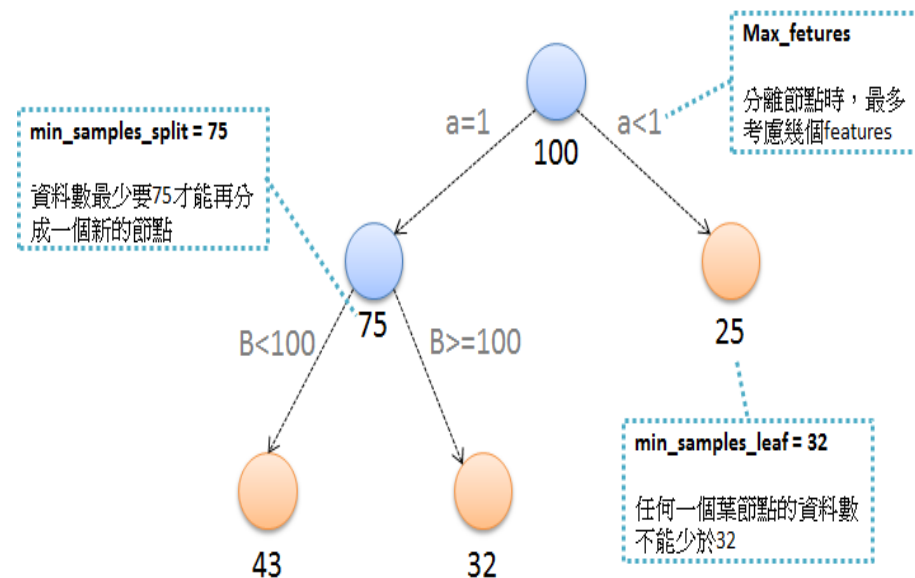
- Smaller imperfect decision tree generally achieves better accuracy.
- Prune leaf nodes recursively to prevent over-fitting.

# Avoid Overfitting in Classification

- The generated tree may **overfit** the training data:
  - Too many branches, some may reflect anomalies due to noise or outliers.
  - Result is in poor accuracy for unseen samples.
- Two approaches to avoid **overfitting**:
  - **Prepruning**: Halt tree construction early—do not split a node if this would result in the goodness measure falling below a threshold.
    - Difficult to choose an appropriate threshold.
  - **Postpruning**: Remove branches from a “fully grown” tree—get a sequence of progressively pruned trees
    - Use a set of data different from the training data to decide which is the “best pruned tree”.

# Prepuning

- Minimum samples for a node split : 資料數目不得小於多少才能再產生新節點。
- Minimum samples for a terminal node (leaf) : 要成為葉節點，最少需要多少資料。
- Maximum depth of tree (vertical depth) : 限制樹的高度最多幾層。
- Maximum number of terminal nodes : 限制最終葉節點的數目
- Maximum features to consider for split : 在分離節點時，最多考慮幾種特徵值。



# Classification And Regression Trees (CART)

# Splitting Based on GINI index

- When a node  $p$  is split into  $k$  partitions (children), the quality of split is computed as,

$$GINI_{split} = \sum_{i=1}^k \frac{n_i}{n} GINI(i)$$

where,  $n_i$  = number of records at child  $i$ ,  
 $n$  = number of records at node  $p$ .

# Attribute Selection using Gini **Index**

- If a data set  $T$  contains examples from  $n$  classes, gini index,  $gini(T)$  is defined as

$$gini(T) = 1 - \sum_{j=1}^n p_j^2$$

where  $p_j$  is the relative frequency of class  $j$  in  $T$ .

- If a data set  $T$  is split into two subsets  $T_1$  and  $T_2$  with sizes  $N_1$  and  $N_2$  respectively, the  $gini$  index of the split data contains examples from  $n$  classes, the  $gini$  index  $gini(T)$  is defined as

$$gini_{split}(T) = \frac{N_1}{N} gini(T_1) + \frac{N_2}{N} gini(T_2)$$

- The attribute provides the smallest  $gini_{split}(T)$  is chosen to split the node (*need to enumerate all possible splitting points for each attribute*).

# Measure of Impurity: GINI

- Gini Index for a given node  $t$  :

$$GINI(t) = 1 - \sum_j [p(j | t)]^2$$

(NOTE:  $p(j | t)$  is the relative frequency of class  $j$  at node  $t$ ).

- Maximum  $(1 - 1/n_c)$  when records are equally distributed among all classes, implying least interesting information
- Minimum (0.0) when all records belong to one class, implying most interesting information

|                   |          |
|-------------------|----------|
| C1                | <b>0</b> |
| C2                | <b>6</b> |
| <b>Gini=0.000</b> |          |

|                   |          |
|-------------------|----------|
| C1                | <b>1</b> |
| C2                | <b>5</b> |
| <b>Gini=0.278</b> |          |

|                   |          |
|-------------------|----------|
| C1                | <b>2</b> |
| C2                | <b>4</b> |
| <b>Gini=0.444</b> |          |

|                   |          |
|-------------------|----------|
| C1                | <b>3</b> |
| C2                | <b>3</b> |
| <b>Gini=0.500</b> |          |

# Examples for Computing GINI

$$GINI(t) = 1 - \sum_j [p(j|t)]^2$$

|    |          |
|----|----------|
| C1 | <b>0</b> |
| C2 | <b>6</b> |

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$Gini = 1 - P(C1)^2 - P(C2)^2 = 1 - 0 - 1 = 0$$

|    |          |
|----|----------|
| C1 | <b>1</b> |
| C2 | <b>5</b> |

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$Gini = 1 - (1/6)^2 - (5/6)^2 = 0.278$$

|    |          |
|----|----------|
| C1 | <b>2</b> |
| C2 | <b>4</b> |

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$Gini = 1 - (2/6)^2 - (4/6)^2 = 0.444$$



# 集成學習

---

# Ensemble Learning

- So far – learning methods that learn a **single hypothesis (model)**, chosen from a hypothesis space that is used to make predictions.


















































---

## “No Free Lunch” Theorem

- No single algorithm wins all the time!
- 

- Ensemble learning
  - select a **collection (*ensemble*) of hypotheses** (models) and **combine their predictions**.
- Example 1 - *generate 100 different decision trees from the same or different training set and have them vote on the best classification for a new example.*

# Example: Weather Forecast

|         |    |    |    |    |    |    |    |
|---------|---|---|---|---|---|---|---|
| 1       |    |    |    |    |    |    |    |
| 2       |    |    |    |    |    |    |    |
| 3       |    |    |    |    |    |    |    |
| 4       |   |   |   |   |   |   |   |
| 5       |  |   |  |  |  |  |  |
| Combine |  |  |  |  |  |  |  |

Combine decisions of multiple models using voting procedure!!!


# Intuition

- Majority vote:  $\epsilon=0.3, n=15 \Rightarrow$   
Binomial Distribution


$$\epsilon_{ensemble} = \sum_{i=8}^{15} \binom{15}{i} \epsilon^i (1-\epsilon)^{15-i} = 0.05$$

- Suppose we have 5 completely independent classifiers...
  - If accuracy is 70% for each
    - $(.7^5) + 5(.7^4)(.3) + 10(.7^3)(.3^2) + \dots$
    - **83.7% majority vote accuracy**
  - 101 such classifiers
    - **99.9% majority vote accuracy**

# Intuition – not any more...



(a) Identical predictive models vs. different predictive models in an ensemble



(b) The different number of predictive models in an ensemble

# Value of Ensembles


- *Key motivation:* reduce the **error rate!!!**
    - *Hope is that it will become much more unlikely that the ensemble will misclassify an example.*
  - Examples: Human ensembles are demonstrably better:
    - *How many jelly beans in the jar?: Individual estimates vs. group average.*
    - *Who Wants to be a Millionaire: Audience vote.*
- 
- *Theory behind:* When combining multiple **independent** and **diverse decisions** each of which is **at least more accurate than random guessing**, random errors cancel each other out, correct decisions are reinforced.

# Achieving Diversity

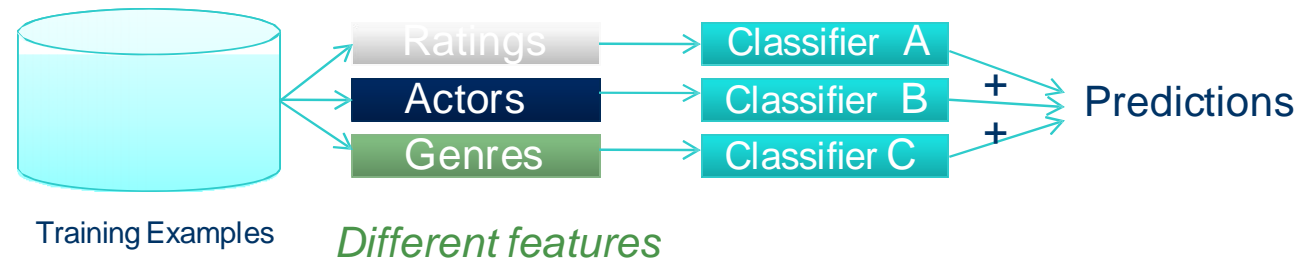
- Using different learning **algorithms**
- Using different hyper-**parameters** in the same algorithm (for example, different polynomial regression)
- Using different input representations, such as using different subsets of input **features**
- Using different training **subsets** of input data, where we have known procedures of **bagging, boosting, and cascading**.

# Achieving Diversity

*Diversity from differences in inputs*



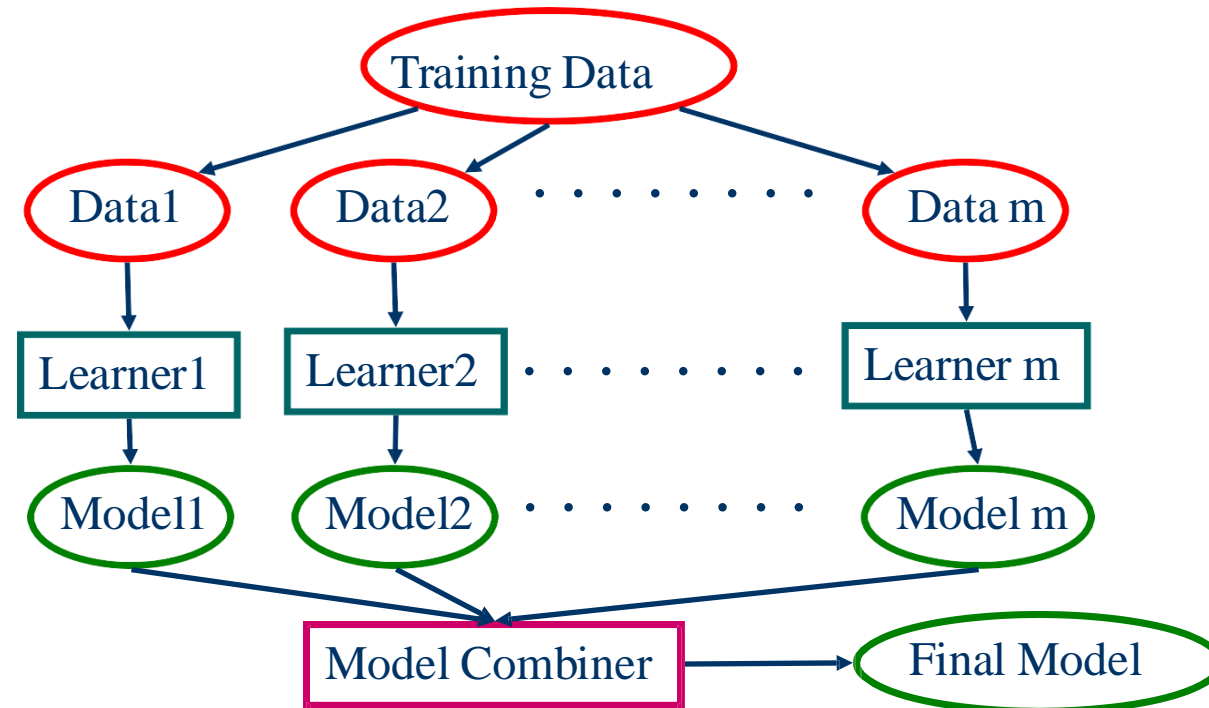
*Divide up training data among models*





# Learning Ensembles

For example: Learn multiple alternative definitions of a concept using: **different training data**

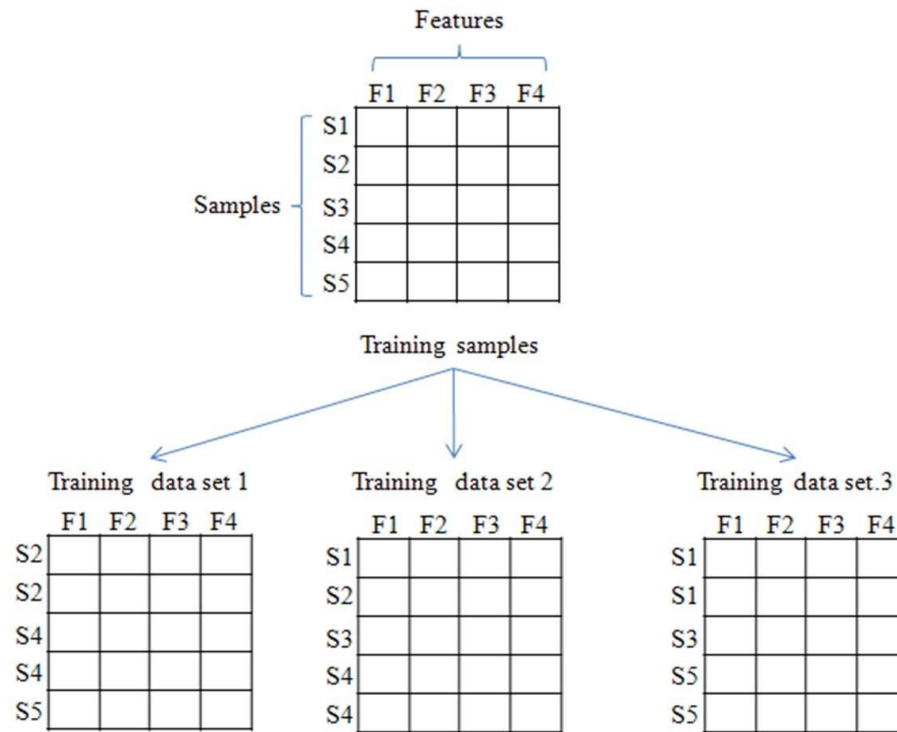


# Homogenous Ensembles



- Use a single, arbitrary learning algorithm but **manipulate training data** to make it learn multiple models.
  - $\text{Data1} \neq \text{Data2} \neq \dots \neq \text{Data } m$
  - $\text{Learner1} = \text{Learner2} = \dots = \text{Learner } m$
- Different methods for changing training data:
  - **Bagging**: Resample training data
  - **Boosting**: Reweight training data

# Bagging



# Bagging

---

- Create ensembles by “*bootstrap aggregation*”, i.e., repeatedly randomly resampling the training data (Brieman, 1996).
  - **Bootstrap**: draw  $N$  items from  $X$  with replacement
- Bagging approach:
  - Train  $M$  learners on  $M$  bootstrap samples
  - Combine outputs by voting (e.g., majority vote)
- Bagging: Decreases error by **decreasing the variance** in the results due to **unstable learners**: algorithms (*like decision trees and neural networks*) whose output can change dramatically when the training data is slightly changed.

# Random Forest

# 隨機森林演算法

- 訓練樣本數為 $N$ ，特徵數為 $M$
- 1. 自訓練樣本以(Bagging, **Bootstrap aggregating**)產生數目 $n(<N)$ 的隨機樣本訓練集
- 2. 對隨機樣本訓練集進行決策樹學習
  - A. 針對每個節點，隨機選取 $m$ 個特徵(取出不放回)
  - B. 依據不同目標函數(entropy, gini, Squared Error)，找出最佳分割
- 3. 重複 $k$ 次(步驟1~2)(訓練出 $k$ 棵決策樹)
- 4. 分類用majority voting；回歸用平均

# Random Forest – practical consideration

- Splits are chosen according to a purity measure:
  - E.g. squared error (regression), Gini index or deviance (classification)
- How to select N?
  - Build trees until the error no longer decreases
- How to select M?
  - Try to recommend defaults, half of them and twice of them and pick the best.

# Adaboost



# Boosting: Strong and Weak Learners



- **Strong Learner** Objective of machine learning
  - Take labeled data for training
  - Produce a classifier which can be *arbitrarily accurate*
- **Weak Learner**
  - Take labeled data for training
  - Produce a classifier which is *more accurate than random guessing*

# Boosting: Strong and Weak Learners

---

- **Weak Learner:** only needs to generate a hypothesis with a training accuracy greater than 0.5, i.e.,  $< 50\%$  error over any distribution
  - *Strong learners are very difficult to construct*
  - *Constructing weaker learners is relatively easy*
- Questions: Can a set of **weak learners** create a single **strong learner** ?

YES 😊


**Boost weak classifiers to a strong learner!!!!** Shapire (1990)

# Boosting: Construct Weak Classifiers




- Idea: Focus on difficult samples which are not correctly classified in the previous steps
- Use different data distribution:
  - Start with **uniform weighting** of samples
  - During each step of learning
    - **Increase weights** of the samples which are **not correctly** learned by the weak learner
    - **Decrease weights** of the samples which are **correctly** learned by the weak learner


# Boosting idea




# Boosting idea




# Boosting idea



# Boosting idea




# Boosting idea





# Boosting idea

Final classifier is  
a combination of weak  
classifiers




# Boosting: Combine Weak Classifiers




- Idea: Better weak classifier gets a larger weight!!!
- Weighted Voting
  - Construct strong classifier by weighted voting of the weak classifiers
  - Weight of each learner is directly proportional to its accuracy

# AdaBoost: Adaptive Boosting

- Each rectangle corresponds to an example, with **weight proportional to its height**.
- Crosses correspond to misclassified examples.
- Size of decision tree indicates **the weight of that hypothesis** in the final ensemble.




## Toy Example




weak classifiers = vertical or horizontal half-planes


## Round 1



## Round 2



### Round 3




$$\epsilon_3 = 0.14$$

$$\alpha_3 = 0.92$$

## Final Classifier


$$H_{\text{final}} = \text{sign} \left( 0.42 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.65 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.92 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} \right)$$





## Final Classifier

$$H_{\text{final}} = \text{sign} \left( 0.42 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.65 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.92 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} \right)$$



*These are principles of **AdaBoost algorithm**: Many variants depending on how to set the weights and how to combine the hypotheses.*

***AdaBoost** € quite popular!!!!*

輸入：訓練集  $D = \{(x_1; y_1), (x_2; y_2), \dots, (x_m; y_m)\}$  ；

基礎學習器  $L$  ； 訓練迴圈輪數  $T$  ；

初始化  $D_1(i) = 1/N$  , 第一次反覆運算時每個樣本的權重都為  $1/N$

For  $t = 1$  to  $T$

在  $D_t$  下訓練： 得到弱假設  $h_t$ ， 即通過訓練集和當前權值分佈得到第  $t$  次的弱學習器  $h_t = L(D, D_t)$

計算  $h_t$  的錯誤率： $\varepsilon_t = \sum D_t(i) [h_t(x_i) \neq y_i]$

令  $\alpha_t = 1/2 \ln((1-\varepsilon_t)/\varepsilon_t)$  此即為  $h_t$  的權重

更新樣本權值：

$Z_t$  為歸一化因數

迴圈結束

輸出：

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} * \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = y_i \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq y_i \end{cases}$$

$$H(X) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(x)\right)$$

# Gradient Boosting

- Gradient Boosting是一種迭代的boosting算法
  - 由Freidman於1999年首次提出
  - 連續訓練多個弱學習器
  - 每次優化前一個弱學習器的損失函數負梯度方向
  - 逐步構建一個強大的集成模型
  - 在每一輪訓練中，Gradient Boosting會計算樣本的損失函數負梯度，用這些負梯度作為新的樣本標籤
  - 並用前一個弱學習器的預測結果作為輸入特徵，進行下一個弱學習器的訓練
  - 每一輪訓練後，Gradient Boosting將新的弱學習器的預測結果與前面所有弱學習器的預測結果相加，從而得到更好的集成模型

# 根據一個人的身高來預測他的體重

1. 初始化模型
  1. 初始化一個弱學習器（例如一個決策樹）
  2. 將其擬合到訓練數據集上。
  3. 這是第一個弱學習器，預測結果可能與真實體重存在誤差
2. 計算殘差
  1. 計算每個樣本的預測誤差，即預測體重和真實體重之間的差異（殘差）
  2. 預測過高的樣本，其殘差為負數，預測過低的樣本，其殘差為正數
3. 訓練下一個弱學習器
  1. 初始化另一個弱學習器，並將其擬合到殘差上
  2. 讓第二個弱學習器學習如何根據第一個弱學習器的預測結果來進行修正
4. 加總預測結果
  1. 連續訓練更多的弱學習器
  2. 每個弱學習器都在前一個弱學習器的殘差上進行優化
  3. 最終將所有弱學習器的預測結果加總起來，得到最終的集成模型

在Gradient Boosting中  
每個弱學習器的預測結果並不是直接相加  
通過乘以學習率（ learning rate ）的值來控制其貢獻

- 預測房價訓練數據：
- 第一個弱學習器的預測結果：1.0,2.0,3.0,4.0,5.0
- 第二個弱學習器的預測結果：0.5,1.5,2.5,3.5,4.5
- 假設學習率（ learning rate ）為0.1
- 將每個弱學習器的預測結果乘以學習率，得到如下：
- 第一個弱學習器加權預測結果：
- 0.1,0.2,0.3,0.4,0.5
- 第二個弱學習器的加權預測結果：
- 0.05,0.15,0.25,0.35,0.45
- 將這兩個弱學習器的加權預測結果相加，得到最終的預測結果：
- 0.15,0.35,0.55,0.75,0.95
- 最終的預測結果 0.15,0.35,0.55,0.75,0.95 就是Gradient Boosting模型對於這五個樣本的預測結果。

# 學習速率 ( learning rate )

- 在Gradient Boosting中
  - 每個弱學習器的學習速率 ( learning rate ) 是一樣的
  - 一個全局的超參數，用來控制每個弱學習器的貢獻程度
  - 學習速率是在算法開始訓練之前由使用者指定在整個訓練過程保持不變。
- 當訓練Gradient Boosting模型時
  - 每個弱學習器的預測結果都會乘以學習速率
  - 然後將所有弱學習器的結果相加得到最終的預測結果
  - 學習速率的值通常是一個介於0和1之間的小數
  - 較小的學習速率意味著每個弱學習器的貢獻變小，模型訓練過程會更加穩定
  - 較大的學習速率可能使模型更快地收斂，但也容易導致過度擬合。
- 訓練Gradient Boosting模型時
  - 使用者需要根據數據集的大小和複雜程度、問題的性質以及計算資源的情況等因素，選擇合適的學習速率值
  - 選擇學習速率是一個重要的過程，通常通過交叉驗證等技術來選擇最佳的學習速率參數，以獲得最好的模型性能

# Gradient Boosting

- Gradient Boosting使用梯度提升演算法
- 邏輯斯回歸或神經網路的訓練採用梯度下降
  - 兩者都是在每一輪反覆運算中利用損失函數相對於模型的負梯度方向的資訊來對當前模型進行更新
  - 在梯度下降中，模型是以參數化形式表示，模型更新等價于參數更新
  - 在梯度提升中，模型不需要進行參數化表示，直接定義在函數空間中擴展了可以使用的模型種類

# Boosting: Summary



- Originally developed by computational learning theorists – Schapire, 1990 (*weak learner*).
- Revised to be a practical algorithm, *AdaBoost*, for building ensembles that empirically improves generalization performance - Freund & Shapire, 1996.

*AdaBoost* - key insights:

- Instead of sampling (as in bagging) re-weight examples!
- Needs smaller number of training samples than bagging
- Final classification based on weighted vote of weak classifiers



# Boosting

---

- Advantages of boosting
  - Flexibility in the choice of weak learners
  - Testing is fast
  - Easy to implement
  - Integrates classification with feature selection
  - Complexity of training may be linear instead of quadratic in the number of training samples
- Disadvantages
  - Often doesn't work for many-class problems

# Ensemble Learning: A Reflection

---

- Combining models adds complexity
  - More difficult to characterize, anticipate predictions, explain predictions, etc.
  - But accuracy may increase
- Violation of Ockham's Razor
  - “simplicity leads to greater accuracy”
  - Identifying the best model requires identifying the proper "model complexity"

# scikit-learn中的boosting方法

- AdaBoost (Adaptive Boosting) :
  - 通過連續訓練多個弱學習器，每次調整樣本權重，專注於前一個弱學習器分類錯誤的樣本
  - 並結合多個弱學習器的預測結果，使得最終預測更準確
- Gradient Boosting :
  - 通過連續訓練多個弱學習器，每次優化前一個弱學習器的損失函數負梯度方向
  - sklearn中的GradientBoostingClassifier用於分類問題，GradientBoostingRegressor用於回歸問題
- XGBoost (Extreme Gradient Boosting) :
  - XGBoost是一種基於Gradient Boosting算法的高效實現
  - 通過樹的並行計算和選擇性的列存儲來加快訓練速度，並具有較好的性能
  - sklearn中的XGBClassifier用於分類問題，XGBRegressor用於回歸問題
- LightGBM :
  - LightGBM是另一種基於Gradient Boosting算法的高效實現
  - 使用基於直方圖的算法來進行分桶 ( binning ) 和加速訓練過程
  - sklearn中的LGBMClassifier用於分類問題，LGBMRegressor用於回歸問題
- CatBoost :
  - CatBoost是俄羅斯Yandex公司開發的一種基於梯度提升算法的高性能機器學習庫
  - sklearn中的CatBoostClassifier用於分類問題，CatBoostRegressor用於回歸問題